

**Xiao Wei** /ˈxiɑʊ weɪ/

Duke University, Durham, NC

[Google Scholar](#) | [Personal Website](#) | [Email](#)

## EDUCATION

**Duke University**, Computer Science, MS

*Durham, NC | Sep. 2024 – Expected May. 2026*

- **GPA:** 3.88 / 4.00
- **Selected Courses:** Robotic Manipulation (*A*), Intro to Medical Robotics in Surgery Tech (*FA25*), Systems for Machine Learning (*A*), Underactuated Robotics (*Self Learning*)

**University of Michigan**, Computer Science and Engineering, BSE

*Ann Arbor, MI | Sep. 2022 – May. 2024*

- **GPA:** 3.97 / 4.00
- **Selected Courses:** Applied Parallel Programming with GPUs (*A*), Continuous Optimization Methods (*A*), Introduction to NLP (*A*), Compiler Construction (*A*), *Game Theory* (*A*)

**UM-SJTU Joint Institute**, Electrical and Computer Engineering, BSE

*Shanghai, China | Sep. 2020 – Aug. 2024*

- **GPA:** 3.80 / 4.00
- **Selected Courses:** Discrete Mathematics (*A*), Honors Mathematics II, III, IV (*A*), Probabilistic Methods in Engineering (*A*), Electronic Circuits (*A+*)

## ACADEMIC PAPERS

- **LAPP: Large Language Model Feedback for Preference-Driven Reinforcement Learning.** Pingcheng Jian, **Xiao Wei**, Yanbaihui Liu, Samuel A. Moore, Michael M. Zavlanos, Boyuan Chen (*under review for TMLR*)
- **dattri: A Library for Efficient Data Attribution.** J Deng, Ting-Wei Li, S Zhang, S Liu, Y Pan, H Huang, X Wang, **X Wei**, P Hu, X Zhang, J Ma. *Neural Information Processing Systems*, volume 37, pages 136763–136781, 2024.

## RESEARCH EXPERIENCE

**Generalizable Learning of Space Division for Hybrid Force-Motion Control** (*Ongoing Research*)

*Project Leader*

*DexLab, Duke University | May. 2025 – Current*

- Learning efficiently from a few demonstrations for the kinematic and dynamical primitives
- Exploring force-motion subspace division automatically with reinforcement learning, based on the primitives
- Proposing to complete table wiping and box flipping tasks with HFMC with online policy

**Large Language Model Feedback for Preference-Driven Reinforcement Learning**

*Core Contributor*

*GRL, Duke University | Sep. 2024 – Mar. 2025*

- Developed LAPP, a novel framework integrating LLM-based preference prediction into reinforcement learning
- Enabled robots to perform complex tasks like gait control and backflips via language-guided behavior tuning
- Achieved superior performance and training efficiency over baselines across multiple benchmarks
- Independently implemented and conducted all experiments for dexterous manipulation tasks

**Efficient Influence Function Calculation through Knowledge Distillation**

*Individual Project*

*TRAIS Lab, UIUC | Jun. 2023 – Jan. 2024*

- Proposed learning data attribution pattern from student model after knowledge distillation
- Reached high attribution similarity in self-distillation and both MLP settings
- Discontinued because of the relatively low performance and weak math guarantee

## ACADEMIC PROJECTS

**Generalizable Compliance Adaptor** (*Ongoing Project*)

*Duke University | Jun. 2025 – Current*

- Designing data augmentation strategies to inject compliance profile into trajectory data and action spaces for policy
- Exploring variable compliance control model from trajectory-only dataset, and collect synthesis data
- Demonstrated a successful whirlwind jump behavior on Unitree GO2 in IsaacGym simulation

**Stage-wise Reward Shaping via Coding LLMs**

*Duke University | Feb. 2025 – May. 2025*

- Automated the design process of stage-wise reward shaping for complex locomotion tasks

- Implemented the pipeline inspired by *Eureka: Human-Level Reward Design via Coding Large Language Models*
- Outperformed manually designed one-phase reward function on “whirlwind jump” locomotion in simulation

#### **dattri: A Library for Efficient Data Attribution**

*TRAIS Lab, UIUC | Feb. 2024 – May. 2024*

- Implemented leave-one-out correlation evaluation method
- Contributed core methods (CG, LiSSA) for influence function calculations
- Built benchmarks on MNIST-10 for logistic regression and MLP models

#### **Performance Recovery of Fully Quantized Models through Adaptors**

*EIC Lab, GaTech | Jun. 2023 – Aug. 2023*

- Proposed to add LoRA to MLP modules (linear layers) of fully quantized vision transformer
- Reached 9.4% increase in ImageNet classification accuracy with efficient fine-tuning

### **TEACHING EXPERIENCE**

---

- **COMPSCI 330: Intro to Design & Analysis of Algorithms** *GTA | Duke University | SP25, FA25*
- **EECS 487: Intro to NLP** *UTA | University of Michigan | FA23*
- **EECS 490: Programming Language** *Grader | University of Michigan | FA23, WN23*

### **HONORS / AWARDS**

---

- **James B. Angell Scholar** *University of Michigan | 2024*
- **The Jackson and Muriel Lum Scholarship** *University of Michigan | 2022*
- **University Honors** *University of Michigan | FA22, WN22, FA23*
- **Dean’s Honor List** *University of Michigan | FA22, WN22, FA23*

### **ACADEMIC SERVICES**

---

- **Journal Reviewer:** IEEE Transactions on Industrial Electronics (TIE)

### **SKILLS**

---

- **Programming Languages:** Python (PyTorch, JAX, Brax), C++, Rust, MATLAB, JavaScript
- **Hardware Experience:** Franka Research 3, Unitree GO2, UR5e
- **Software Experience:** MuJoCo, IssacGym, IsaacLab, Deoxys, Franky